

Philadelphia Crime

Nicholas Morris

Machine Learning

Wrangling

Weekly Crime Rate

Computing weekly crime rate by district. Adding previous weeks of crime to the data.

Feature Engineering

Additional Data


Adding economic trends such as unemployment and CPI. Extracting timestamp components such as year and quarter. Converting timestamp components and district to binary data points.

Modeling

Predictions

Training linear regression, XGBoost, and deep learning neural network models. Evaluating performance. Computing feature drift to signal retraining.

Wrangling



—

Dataset

Below is the first and last arrest in the data. There are 2,237,605 total arrests. The key columns are Dispatch_Date and Dc_Dist which tell us when and where the arrests occur. [\[Link to the dataset\]](#)

Dc_Dist	Psa	Dispatch_Date_Time	Dispatch_Date	Dispatch_Time	Hour	Dc_Key	Location_Block	UCR_General	Text_General_Code	Police_Districts	Month	Lon	Lat
3	I	2006-01-01 00:00:00	2006-01-01	00:00:00	0	200603000002	S 8TH ST /SOUTH ST	2600.0	All Other Offenses	3.0	2006-01	-75.155491	39.942416
24	3	2017-03-23 01:29:00	2017-03-23	01:29:00	1	201724026395	3700 BLOCK RICHMOND ST	400.0	Aggravated Assault No Firearm	17.0	2017-03	-75.087735	39.99009

Crime Rate

Dispatch_Date	Dc_Dist	Crimes
---------------	---------	--------

2006-01-01	1	14.0
------------	---	------

2006-01-02	1	17.0
------------	---	------

2017-03-22	92	0.0
------------	----	-----

2017-03-23	92	0.0
------------	----	-----


Aggregating the data by
Dispatch_Date and Dc_Dist we get
102,500 rows of crime data.

Previous Weeks Of Crime

Aggregating the data by Year, Week, and Dc_Dist we get 14,575 rows of crime data. For each district the previous four weeks of crime were extracted.

Dc_Dist	Year	Week	Crimes	Crimes(t-1)	Crimes(t-2)	Crimes(t-3)	Crimes(t-4)
1	2006	4	97.0	144.0	103.0	129.0	14.0
1	2006	5	110.0	97.0	144.0	103.0	129.0
1	2006	6	77.0	110.0	97.0	144.0	103.0
1	2006	7	83.0	77.0	110.0	97.0	144.0

Feature Engineering



—

Time Components

Year	Quarter	Month	Week	Days	Crimes
2006	1	1	4	7	97.0
2006	1	1	5	7	110.0
2006	1	2	6	7	77.0
2006	1	2	7	7	83.0

The Year, Quarter, Month, Week, and number of Days in a week were extracted from Dispatch_Date.

Economic Data

The NASDAQ closing price, Unemployment rate, CPI, PPI, GDP, GDI, and Federal Funds Rate were pulled from FRED (Federal Reserve Economic Data).


Year	Week	NASDAQ	Unemployment	CPI	PPI	GDP	GDI	Federal_Funds_Rate
2006	4	2281.437143	4.700000	0.762195	164.300000	13599.16	13795.893	4.290000
2006	5	2284.638571	4.771429	0.361852	162.514286	13599.16	13795.893	4.432857
2006	6	2258.892857	4.800000	0.201715	161.800000	13599.16	13795.893	4.490000
2006	7	2274.302857	4.800000	0.201715	161.800000	13599.16	13795.893	4.490000

The time components and Dc_Dist were converted to binary variables.

Binary Data

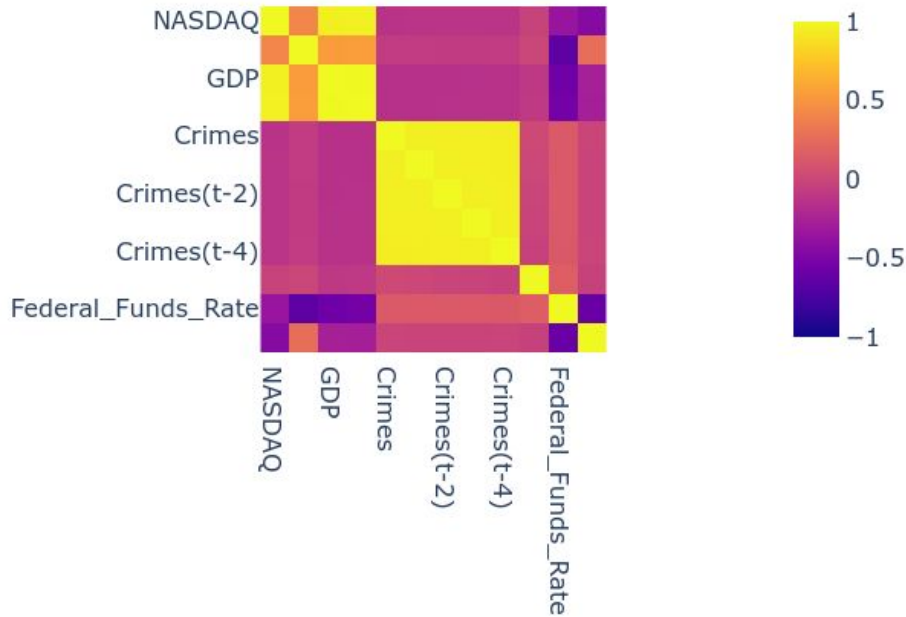
Year_2006	...	Week_4	Week_5	...	Dc_Dist_1	...	Dc_Dist_92	Crimes
1	...	1	0	...	1	...	0	97.0
1	...	0	1	...	1	...	0	110.0
1	...	0	0	...	1	...	0	77.0
1	...	0	0	...	1	...	0	83.0

Data Exploration



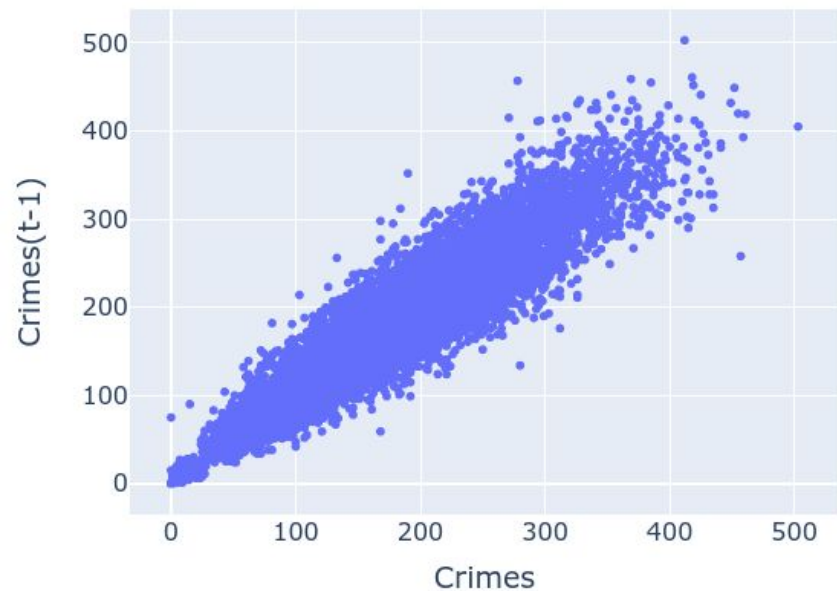
—

Correlation Heatmap




There's two hot spots for relationships in the data. The top-left yellow region shows strong relationships between the economic indicators. The centered yellow region shows strong relationships between crime and previous weeks of crime. There's also a strong relationship in the bottom-left where it is dark purple for the federal funds rate and the other economic indicators.

Crimes vs. Crimes(t-1)




We can see that the previous week of crime is similar to the current week of crime. When the previous week is high or low the current week follows the same trend.

Crimes vs. Dc_Dist




We can see that crime is higher in some districts and lower in others.

Crimes vs. Month




We can see that crime is slightly higher in the summer than the winter. In August, the median crime rate is 173 arrests per week. In January, the median crime rate is 135 arrests per week. That's a 22% decrease.

Crimes vs. Year



We can see that crime is slowly decreasing over time. In 2006, the median crime rate is 176 arrests per week. Ten years later in 2016, the median crime rate is 137.5 arrests per week. That's a 22% decrease.

Modeling



—

Model Parameters

Linear Regression

Library: scikit-learn
Length Of Path: 1e-9
Number Of Alphas: 16
Cross Validation Folds: 3
Tolerance: 1e-4
Max Iterations: 500

XGBoost

Library: xgboost
Boosting Rounds: 100
Learning Rate: 0.1
Max Depth: 7
Min Child Weight: 1
Column Sampling: 0.8
Row Sampling: 0.8

Neural Network

Library: Tensorflow
Epochs: 500
Learning Rate: 0.001
Batch Size: 16
Layers: 10
Nodes Per Layer: 128
Solver: Adam

Model Comparison

Linear Regression

R2: 0.94

RMSE: 20.7

In Control: 97.39%

Model Indicators:

1. Crimes(t-1)
2. Crimes(t-2)
3. Crimes(t-3)
4. Crimes(t-4)
5. Week_52

XGBoost

R2: 0.93

RMSE: 21.9

In Control: 97.6%

Model Indicators:

1. Crimes(t-2)
2. Crimes(t-3)
3. Crimes(t-1)
4. Crimes(t-4)
5. Dc_Dist_15

Neural Network

R2: 0.89


RMSE: 28.4

In Control: 97.36%

Model Indicators:


1. GDI
2. Federal_Funds_Rate
3. GDP
4. PPI
5. Unemployment

Parity Plot



These predictions come from the linear regression model. These predictions are done on 20% of the data that the model did not see during training. The model trained on data from 2006, week 4 to 2015, week 1. These predictions are done on data from 2015, week 1 to 2017, week 12. The predictions are centered on the red line (perfect predictions). There is a tendency to over-predict the crime rate.


Histogram For Residuals, 97.39% In Control



The residuals are prediction error = actual - predicted.

The residuals have a tight bell shape, which is good, and they are centered on zero. Control limits were computed on the residuals and we can see that the prediction error is mostly under control. We can see a skew to the left, which shows the model has a tendency to over-predict the crime rate.


Histogram For R2



The prediction error was resampled 1000 times at a 50% sampling rate with replacement. Then R2 was computed on each sample to get a distribution.

R2 has a tight range between 0.92 and 0.95, which is good. R2 has a bell shape, which is good, and a slight skew to the left.


Histogram For RMSE



The prediction error was resampled as previously mentioned to get a distribution for RMSE.


On average, the predictions are off by 19 to 23 crimes per week, which is a tight range.

Feature Drift, Drift Detected If $pvalue < 0.05$



A Kolmogorov-Smirnov test was performed for each column in the data to see if the distribution of the testing data is the same as the training data. If the testing data does not share the same distribution as the training data, then there is a drift, which signals for model retraining. Most of the columns do not experience a drift, which is good; but some of them do.


Feature Drift, Drift Detected If pvalue < 0.05



We can see that the crime rate, the years, and the economic data are experiencing a drift. This is because crime is decreasing, the years are different between the training and testing data, and most of the economic data has a trend to increase.

The linear regression model was retrained to include the testing data.

Deployment




—

The data we start with.

The latest data we
want predictions for.

Retrain the model on
the initial data and
new data.



Data wrangling,
feature engineering,
model training.

See if the distribution
of the new data is
significantly different
than the initial data.



Thank You